

ECS7001P - NEURAL NETWORKS AND NLP

School of Electronic Engineering and Computer Science, Queen Mary University of London, UK

Assignment 1: Embeddings, Text Classification, And Machine Translation

SHIVANI GURUNG RAKESH: 210268155

A. Word Embeddings with Word2Vec

1. Preprocessing the training corpus

```
4 [['[', 'Sense', 'and', 'Sensibility', 'by', 'Jane',  
  'Austen', '1811', '']], ['CHAPTER', '1'], ['But', ' ',  
  'then', ' ', 'if', 'Mrs', '.', 'Dashwood', 'should', 'live',  
  'fifteen', 'years', 'we', 'shall', 'be', 'completely',  
  'taken', 'in', '."'], ['"', 'Fifteen', 'years', '!']]  
  
2 ['sense sensibility jane austen', 'mrs dashwood live  
  fifteen years shall completely taken']
```

2. Creating the corpus vocabulary and preparing the dataset

Number of unique words: 10098

```
Sample word2idx: [('sense', 0), ('sensibility', 1),  
  ('jane', 2), ('austen', 3), ('family', 4), ('dashwood', 5),  
  ('long', 6), ('settled', 7), ('sussex', 8), ('estate', 9)]
```

```
Sample idx2word: [(0, 'sense'), (1, 'sensibility'), (2,  
  'jane'), (3, 'austen'), (4, 'family'), (5, 'dashwood'), (6,  
  'long'), (7, 'settled'), (8, 'sussex'), (9, 'estate')]
```

```
Sample sents_as_id: [[0, 1, 2, 3], [68, 5, 194, 592, 33,  
  593, 285, 594]]
```

3. Building the skip-gram neural network architecture

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 1)]	0	[]
input_2 (InputLayer)	[(None, 1)]	0	[]
target_embed_layer (Embedding)	(None, 1, 100)	1009800	['input_1[0][0]']
context_embed_layer (Embedding)	(None, 1, 100)	1009800	['input_2[0][0]']
reshape (Reshape)	(None, 100)	0	['target_embed_layer[0][0]']
reshape_1 (Reshape)	(None, 100)	0	['context_embed_layer[0][0]']
dot (Dot)	(None, 1)	0	['reshape[0][0]', 'reshape_1[0][0]']
dense (Dense)	(None, 1)	2	['dot[0][0]']

=====
Total params: 2,019,602
Trainable params: 2,019,602
Non-trainable params: 0
=====

4. Training the models (and reading McCormick's tutorial)

- a. What would the inputs and outputs to the model be?

The input would be a numeric/vector representation of textual data/word/string. The outcome is a single vector representation for each word that explains us how probable it is to be picked as the next word.

- b. How would you use the Keras framework to create this architecture?

This layer takes as inputs a target word and a context word. The embedding in the preceding layer, as well as the embedding's change in the reshaping layer. Finally, the dot product is taken into account.

- c. What are the reasons this training approach is considered inefficient?

This training approach is considered inefficient because Internally the terms provided here do not sufficiently represent semantic connections since they do not capture the contextual meaning. A huge dataset, such as the one created artificially for word2vec, is required.

5. Getting the word embeddings

```
from pandas import DataFrame

print(DataFrame(word_embeddings, index=idx2word.values()).head(10))
```

	0	1	2	3	4	5	\
sense	-0.004377	0.011746	-0.012830	-0.009580	-0.020550	-0.014081	
sensibility	0.003382	0.000622	0.022603	0.030473	0.017281	0.023543	
jane	0.089320	0.010615	-0.082091	0.033990	0.015080	0.165932	
austen	-0.026869	-0.021149	0.016017	0.043684	-0.000899	-0.000202	
family	0.004196	-0.017450	-0.034747	0.027821	0.072123	0.012793	
dashwood	-0.077659	0.020736	-0.064296	-0.015214	-0.024381	0.064664	
long	0.080348	0.024557	0.025339	-0.112085	0.104181	0.079189	
settled	-0.044766	0.028073	0.087832	-0.096838	0.034346	-0.003181	
sussex	-0.028548	0.002718	0.023980	-0.012256	0.014442	0.012626	
estate	-0.037755	-0.041832	-0.022449	0.004864	-0.003568	0.023726	


	6	7	8	9	...	90	91	\
sense	-0.019290	-0.005767	-0.006455	0.002863	...	0.009058	0.004660	
sensibility	-0.014133	0.044463	-0.008119	-0.000441	...	-0.000632	0.015275	
jane	-0.010550	-0.045091	0.055103	0.051547	...	0.073622	-0.043846	
austen	0.005886	-0.034634	-0.003676	0.015691	...	-0.012423	0.008044	
family	-0.048012	0.059493	-0.014547	-0.007315	...	0.097147	0.085674	
dashwood	-0.100502	-0.082606	0.022895	0.081411	...	-0.076113	-0.160559	
long	-0.009862	0.161314	0.021617	0.029898	...	-0.028522	0.032630	
settled	-0.006206	-0.041104	-0.005688	-0.060299	...	-0.051511	0.059097	
sussex	-0.015412	-0.013735	-0.003959	-0.029497	...	-0.001287	0.012682	
estate	0.025047	0.035452	-0.023745	0.011232	...	-0.028287	-0.001712	

	92	93	94	95	96	97	\
sense	0.022175	0.019816	-0.017966	-0.011483	0.018054	-0.017310	
sensibility	-0.042384	0.012944	-0.001387	0.004599	0.022433	-0.023866	
jane	0.013386	-0.004306	-0.072449	-0.089178	0.035184	-0.076658	
austen	0.021220	-0.009510	0.001668	0.006835	0.036758	-0.010606	
family	-0.056270	0.012475	0.028407	-0.058864	-0.008543	0.054817	
dashwood	-0.084238	0.078736	0.068596	-0.108069	0.096946	0.066056	
long	0.017362	0.036617	-0.048536	0.082366	-0.083572	-0.040926	
settled	-0.009625	-0.044698	0.009037	-0.088269	0.006469	-0.040118	
sussex	-0.013260	0.038690	-0.005735	0.000106	0.002469	0.022721	
estate	-0.035640	0.028941	0.042068	0.095949	-0.038173	0.049339	

	98	99
sense	0.000966	0.012991
sensibility	-0.032437	-0.023564
jane	-0.036430	0.028108
austen	-0.005499	0.028972
family	-0.053529	-0.036122
dashwood	-0.042492	-0.012248
long	0.014095	0.047280
settled	0.038472	0.049939
sussex	0.021741	0.009505
estate	-0.059784	-0.018636

[10 rows x 100 columns]

6. Exploring and visualizing your word embeddings using t-SNE



B. Using LSTMs for Text Classification

1. Section 2, reading the inputs for the LSTM

```
print('Length of sample train_data before preprocessing:', len(train_data[0]))
print('Length of sample train_data after preprocessing:', len(padded_train_data[0]))
print('Sample train data:', padded_train_data[0])
```

[illegible]

2. Building the model

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 500)]	0
embed_layer (Embedding)	(None, 500, 100)	1000000
lstm_1 (LSTM)	(None, 100)	80400
output_layer (Dense)	(None, 1)	101


=====
Total params: 1,080,501
Trainable params: 1,080,501
Non-trainable params: 0

3. Section 4, training the model

```
history = model.fit(train_x, train_y, epochs=6, batch_size=1000, validation_data=(validation_x, validation_y))
```

```
Epoch 1/6
23/23 [=====] - 37s 1s/step - loss: 0.6265 - accuracy: 0.6451 - val_loss: 0.4710 - val_accuracy: 0.7910
Epoch 2/6
23/23 [=====] - 22s 949ms/step - loss: 0.3790 - accuracy: 0.8531 - val_loss: 0.3176 - val_accuracy: 0.8715
Epoch 3/6
23/23 [=====] - 22s 955ms/step - loss: 0.2510 - accuracy: 0.9057 - val_loss: 0.2899 - val_accuracy: 0.8875
Epoch 4/6
23/23 [=====] - 24s 1s/step - loss: 0.1957 - accuracy: 0.9307 - val_loss: 0.2733 - val_accuracy: 0.8960
Epoch 5/6
23/23 [=====] - 22s 954ms/step - loss: 0.1563 - accuracy: 0.9486 - val_loss: 0.2805 - val_accuracy: 0.8905
Epoch 6/6
23/23 [=====] - 22s 957ms/step - loss: 0.1289 - accuracy: 0.9593 - val_loss: 0.3006 - val_accuracy: 0.8945
```

```
plot_history(history.history, path="standard.png")
plt.show()
```



Based on the accuracy plot, what do you think the optimal stopping point for your model should have been?

```
[ ] len(test_data)
```

25000

4. Evaluating the model on the test data

Evaluate the model on the padded test data using the code in the following cell block.

```
[ ] # YOUR CODE TO EVALUATE THE MODEL ON TEST DATA GOES HERE
results = model.evaluate(padded_test_data, test_labels)
print('test_loss:', results[0], 'test_accuracy:', results[1])

782/782 [=====] - 110s 139ms/step - loss: 0.3601 - accuracy: 0.8685
test_loss: 0.3600790500640869 test_accuracy: 0.8684800267219543
```

5. Section 6, extracting the word embedding

Sanity Check

Print the shape of the word embeddings using the line of code below. It should return (VOCAB_SIZE, EMBED_SIZE)

```
[ ] print('Shape of word_embeddings:', word_embeddings.shape)
```

Shape of word_embeddings: (10000, 100)

6. Visualizing the reviews

```
print(' '.join(idx2word[idx] for idx in train_data[0]))
```

<START> this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert <UNK> is an amazing actor and now the same being director <UNK> father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for <UNK> and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also <UNK> to the two little boy's that played the <UNK> of norman and paul they were just brilliant children are often left out of the <UNK> list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all

7. Visualizing the word embeddings

```
from pandas import DataFrame

print(DataFrame(word_embeddings, index=idx2word.values()).head(10))
```

	0	1	2	3	4	5	\
woods	0.021322	-0.016948	0.016021	-0.001831	0.003546	0.018424	
hanging	-0.028856	0.000362	-0.005027	0.004653	0.011076	-0.019683	
woody	-0.002695	0.018942	-0.016365	-0.008425	0.020334	-0.023362	
arranged	-0.014585	-0.000727	-0.016299	-0.011049	0.006553	0.013894	
bringing	0.008422	0.010303	-0.001404	0.012576	0.019186	-0.013220	
wooden	0.011720	-0.005625	0.011642	-0.028148	0.012258	-0.014192	
errors	0.025149	-0.017638	0.017255	-0.006253	-0.015176	0.020370	
dialogs	-0.026621	0.008397	-0.018563	-0.003514	-0.014060	0.020323	
kids	0.011181	0.007972	0.001516	0.006886	-0.004066	0.011590	
uplifting	-0.012604	-0.005350	-0.017992	-0.005206	-0.020141	0.019626	

	6	7	8	9	...	90	91	\
woods	0.016331	-0.010687	0.009252	0.009148	...	-0.010355	0.011949	
hanging	-0.023211	-0.002756	0.008896	0.009968	...	0.019486	-0.017137	
woody	0.015943	-0.009274	0.007285	0.009084	...	0.020612	-0.005199	
arranged	-0.023375	-0.009402	-0.018564	-0.021615	...	0.011507	0.013424	
bringing	-0.015699	0.002351	0.001682	-0.003326	...	0.019078	-0.006840	
wooden	0.021057	-0.024747	-0.002334	-0.011708	...	-0.017786	0.016806	
errors	-0.013985	-0.004677	-0.005856	0.010407	...	-0.000786	-0.013370	
dialogs	0.009236	0.013456	0.016642	0.008468	...	0.002678	-0.020391	
kids	-0.019377	-0.005905	0.027780	0.002733	...	0.016922	0.006767	
uplifting	-0.002850	0.009913	0.015541	-0.012341	...	0.013410	-0.015321	

```

          92      93      94      95      96      97  \
woods      -0.001569 -0.015296  0.016905 -0.001719 -0.018573 -0.016062
hanging     0.017108 -0.015979  0.012520  0.012075  0.020443 -0.008092
woody       -0.025347 -0.002880  0.006627 -0.005779 -0.005327  0.021822
arranged    0.003250  0.020580 -0.019248 -0.013976  0.009721 -0.017364
bringing   -0.000961  0.005473  0.029043 -0.019207  0.028355  0.000540
wooden      0.016476 -0.018995  0.023646  0.005147  0.009694  0.003889
errors     -0.015242  0.011041 -0.020507  0.014845  0.016741 -0.005231
dialogs     -0.006387 -0.011488 -0.001304  0.000383 -0.012729 -0.006082
kids        -0.002744 -0.019494  0.006983 -0.016619  0.022346  0.013522
uplifting   -0.003392  0.004667  0.018703 -0.008488  0.023051 -0.018911

          98      99
woods      0.003212  0.010031
hanging     0.003593 -0.006954
woody       0.018661  0.006990
arranged    -0.022717  0.012371
bringing    -0.001960  0.021795
wooden      -0.005483  0.019206
errors      -0.010556 -0.011302
dialogs     0.014426 -0.007965
kids        0.008722  0.010755
uplifting   -0.018329  0.010686

[10 rows x 100 columns]
```

8.

Section 9

- a. Create a new model that is a copy of model step 3. To this new model, add two dropout layers, one between the embedding layer and the LSTM layer and another between the LSTM layer and the output layer. Repeat steps 4 and 5 for this model. What do you observe?


```
[ ] new_model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 500)]	0
embed_layer (Embedding)	(None, 500, 100)	1000000
dropout_layer_1 (Dropout)	(None, 500, 100)	0
lstm_1 (LSTM)	(None, 100)	80400
dropout_layer_2 (Dropout)	(None, 100)	0
output_layer (Dense)	(None, 1)	101

=====
Total params: 1,080,501
Trainable params: 1,080,501
Non-trainable params: 0

```
plot_history(new_history.history, path="standard.png")  
plt.show()
```



The graphs before and after we applied dropout to the model may be seen. Even after running the whole epochs, we can plainly observe that after adding the dropout, our validation loss does not rise. This demonstrates that our model is no longer over fitting. We can also see that after we introduced the dropouts, the validation accuracy has improved somewhat. We might be able to improve accuracy and reduce over fitting by increasing dropout rates even further.


C. Comparing The Classification Models

1. Build a neural network classifier using one-hot word vectors (Model 1), and train and evaluate it

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256)]	0
lambda (Lambda)	(None, 256, 10000)	0
global_average_pooling1d_masked (GlobalAveragePooling1DMasked)	(None, 10000)	0
dense (Dense)	(None, 16)	160016
dense_1 (Dense)	(None, 1)	17

=====
Total params: 160,033
Trainable params: 160,033
Non-trainable params: 0
=====

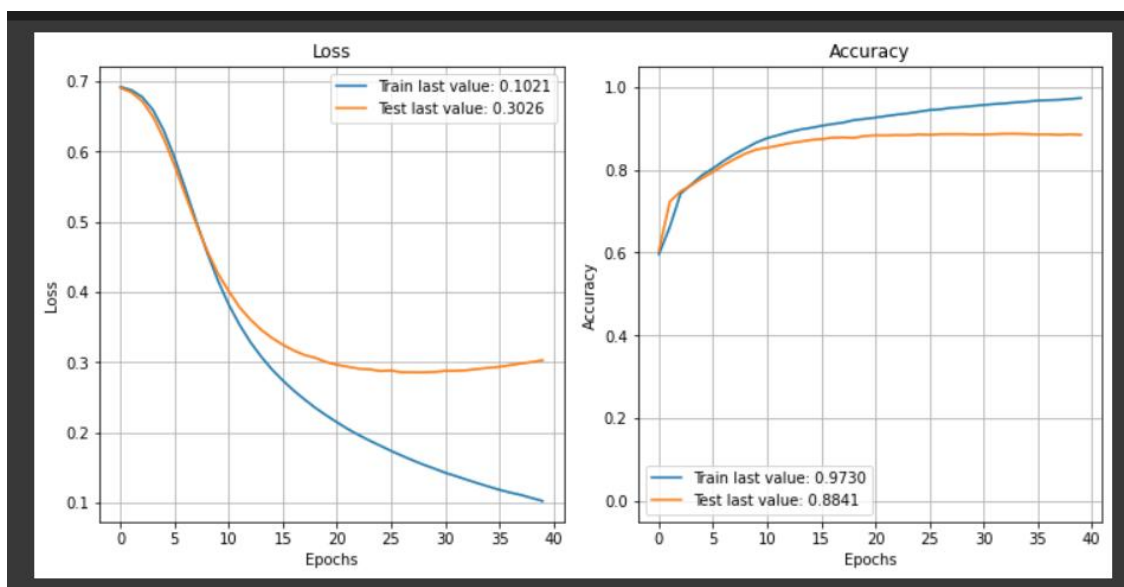


2. Modify your model to use a word embedding layer instead of one-hot vectors (Model 2), and to learn the values of these word embedding vectors along with the model

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 256)]	0
embedding (Embedding)	(None, 256, 16)	160000
global_average_pooling1d_masked_1 (GlobalAveragePooling1DMasked)	(None, 16)	0
dense_2 (Dense)	(None, 16)	272
dense_3 (Dense)	(None, 1)	17

=====
Total params: 160,289
Trainable params: 160,289
Non-trainable params: 0
=====




- Adapt your model to load and use pre-trained word embeddings instead (Model 3); train and evaluate it and compare the effect of freezing and fine-tuning the embeddings

Model: "model_2"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 256)]	0
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_2 (GlobalAveragePooling1DMasked)	(None, 300)	0
dense_4 (Dense)	(None, 16)	4816
dense_5 (Dense)	(None, 1)	17

=====
Total params: 120,005,133
Trainable params: 4,833
Non-trainable params: 120,000,300
=====

```
plot_history(history_3.history, path="standard.png")  
plt.show()
```



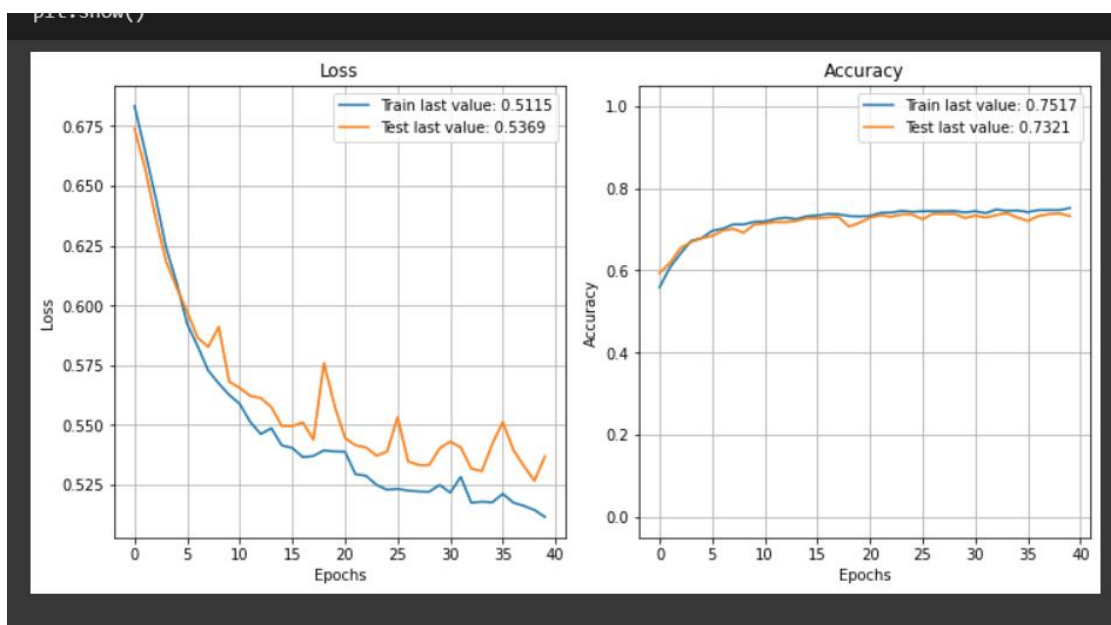
- One way to improve the performance is to add another fully-connected layer to your network. Try this (Model 4) and see if it improves the performance. If not, what can you do to improve it?

Model: "model_5"


Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, 256)]	0
Glove_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_4 (GlobalAveragePooling1D)	(None, 300)	0
dense_9 (Dense)	(None, 100)	30100
dense_10 (Dense)	(None, 16)	1616
dense_11 (Dense)	(None, 1)	17

=====
Total params: 120,032,033
Trainable params: 31,733
Non-trainable params: 120,000,300
=====

- Training and validation loss plot after adding one dense layer



- Training and validation loss plot after adding two dense layer



As layers are added, the number of weights in the network increases, lowering the model's complexity. Without a large training set, an increasingly huge network is prone to overfit, lowering accuracy on test data. The two plots show that test accuracy is decreasing while train accuracy is increasing, signalling that our model is overfitting. We can improve performance by adding dropouts, adjusting hyperparameters, and even simplifying the model.

5. Build a CNN classifier (Model 5), and train and evaluate it. Then try adding extra convolutional layers, and conduct training and evaluation.
 - one convolutional layer


```
Model: "model_7"
```

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 256)]	0
embed_layer (Embedding)	(None, 256, 300)	3000000
conv1d (Conv1D)	(None, 251, 100)	180100
global_max_pooling1d (GlobalMaxPooling1D)	(None, 100)	0
dense_16 (Dense)	(None, 1)	101

```

Total params: 3,180,201
Trainable params: 3,180,201
Non-trainable params: 0

```




- Extra convolutional layer

```

Model: "model_8"
Layer (type)                 Output Shape              Param #
=====
input_9 (InputLayer)         [(None, 256)]             0
embed_layer (Embedding)      (None, 256, 300)         3000000
conv1d_1 (Conv1D)            (None, 251, 100)         180100
conv1d_2 (Conv1D)            (None, 246, 100)         60100
global_max_pooling1d_1 (Glo (None, 100)              0
balMaxPooling1D)
dense_17 (Dense)             (None, 1)                101
=====
Total params: 3,240,301
Trainable params: 3,240,301
Non-trainable params: 0

```

The validation accuracy has decreased and the validation loss has increased as our model has begun to overfit. The model's complexity rises as the layered complexity rises, which is presumably why it's overfit.

D. A Real Text Classification Task

1. Preprocess the data, to adapt the models from Parts C - Done in Jupiter notebook
2. Adapt your models without pre-trained word embeddings in Part C to this task

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 128)]	0
embedding (Embedding)	(None, 128, 100)	789800
global_average_pooling1d (GlobalAveragePooling1D)	(None, 100)	0
dense (Dense)	(None, 16)	1616
dense_1 (Dense)	(None, 3)	51

=====
 Total params: 791,467
 Trainable params: 791,467
 Non-trainable params: 0

The accuracy of the validation data was 55.9%. The accuracy of the test data was 58.6%.

3. Adapt your models with pre-trained word embeddings in Part C to this task (Model 2); train and evaluate it

Neural bag of words using pre-trained word embeddings

Model: "model_3"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 128)]	0
GloVe_Embeddings (Embedding)	(None, 128, 300)	120000300
global_average_pooling1d_2 (GlobalAveragePooling1D)	(None, 300)	0
dense_6 (Dense)	(None, 16)	4816
dense_7 (Dense)	(None, 3)	51

=====
Total params: 120,005,167
Trainable params: 4,867
Non-trainable params: 120,000,300

4. Build and evaluate two more classifiers with multiple inputs (Model 3: separate inputs for text and aspect)

Model: "model_6"

Layer (type)	Output Shape	Param #	Connected to
input_layer_6 (InputLayer)	[(None, 16)]	0	[]
input_layer_7 (InputLayer)	[(None, 128)]	0	[]
GloVe_Embeddings (Embedding)	multiple	120000300	['input_layer_6[0][0]', 'input_layer_7[0][0]']
global_average_pooling1d_4 (GlobalAveragePooling1D)	(None, 300)	0	['GloVe_Embeddings[3][0]']
global_average_pooling1d_5 (GlobalAveragePooling1D)	(None, 300)	0	['GloVe_Embeddings[4][0]']
dense_12 (Dense)	(None, 16)	4816	['global_average_pooling1d_4[0][0]']
dense_13 (Dense)	(None, 16)	4816	['global_average_pooling1d_5[0][0]']
concatenate (Concatenate)	(None, 32)	0	['dense_12[0][0]', 'dense_13[0][0]']
dense_14 (Dense)	(None, 3)	99	['concatenate[0][0]']

=====
Total params: 120,010,031
Trainable params: 9,731
Non-trainable params: 120,000,300

5. Build and evaluate the classifier extracting information from LSTM (Model 4)

Model: "model_8"			
Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[(None, 128)]	0	[]
GloVe_Embeddings (Embedding)	multiple	120000300	['input_7[0][0]']
bidirectional (Bidirectional)	(None, 128, 200)	320000	['GloVe_Embeddings[7][0]']
input_8 (InputLayer)	[(None, 128)]	0	[]
dot (Dot)	(None, 200)	0	['bidirectional[0][0]', 'input_8[0][0]']
dense_18 (Dense)	(None, 16)	3216	['dot[0][0]']
dense_19 (Dense)	(None, 3)	51	['dense_18[0][0]']
=====			
Total params: 120,324,367			
Trainable params: 324,067			
Non-trainable params: 120,000,300			

E. Neural Machine Translation

1. Implementing the encoder

```
"""
Task 1 encoder

Start
"""
# The train encoder
# (a.) Create two randomly initialized embedding lookups, one for the source, another for the target.
print('Task 1(a): Creating the embedding lookups...')
embeddings_source = Embedding(input_dim=self.vocab_source_size, output_dim=self.embedding_size, mask_zero=True, trainable=True)
embeddings_target = Embedding(input_dim=self.vocab_target_size, output_dim=self.embedding_size, mask_zero=True, trainable=True)

# (b.) Look up the embeddings for source words and for target words. Apply dropout to each encoded input
print('\nTask 1(b): Looking up source and target words...')
source_word_embeddings = Dropout(self.embedding_dropout_rate)(embeddings_source(source_words))
target_words_embeddings = Dropout(self.embedding_dropout_rate)(embeddings_target(target_words))

# (c.) An encoder LSTM() with return sequences set to True
print('\nTask 1(c): Creating an encoder')

encoder_outputs, encoder_state_h, encoder_state_c = LSTM(self.hidden_size, return_sequences=True, return_state=True)(source_word_embeddings)

"""
End Task 1
"""
```

We will begin by creating the embedding layers for the source and target. A random initialization will be used to start the embeddings, and they will be taught during training. The input dimensions will be the source and target dimensions. They will both be the same size for embedding and output. We will set "mask zero" to true because we wish to remove the padding. Next, dropouts will be added to the embeddings for both the source and destination. An estimation of the dropout rate has already been made.

The inputs will be sent to the embedding layers, who will then pass them on to the dropout layers. Finally, we'll create an LSTM layer to handle the output of the dropout layers. To get the outputs, set "return sequence" to true, and "return state" to true to access the LSTM's hidden states.

2. In this step we will be implementing the decoder

```

"""
Task 2 decoder for inference

Start
"""
# Task 1 (a.) Get the decoded outputs
print('\n Putting together the decoder states')
# get the initial states for the decoder, decoder_states
# decoder states are the hidden and cell states from the training stage

decoder_states = [decoder_state_input_h, decoder_state_input_c]

# use decoder states as input to the decoder lstm to get the decoder outputs, h, and c for test time inference
decoder_outputs_test, decoder_state_output_h, decoder_state_output_c = decoder_lstm(target_words_embeddings, initial_state=decoder_states)

# Task 1 (b.) Add attention if attention
if self.use_attention:
    decoder_attention = AttentionLayer()
    decoder_outputs_test = decoder_attention([encoder_outputs_input, decoder_outputs_test])

# Task 1 (c.) pass the decoder_outputs_test (with or without attention) to the decoder dense layer
decoder_outputs_test = decoder_dense(decoder_outputs_test)

"""
End Task 2
"""

```

- Decoder interface model summary

Putting together the decoder states

Decoder Inference Model summary			
Model: "model_2"			
Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, None)]	0	[]
embedding_1 (Embedding)	(None, None, 100)	250600	['input_2[0][0]']
dropout_1 (Dropout)	(None, None, 100)	0	['embedding_1[0][0]']
input_3 (InputLayer)	[(None, 200)]	0	[]
input_4 (InputLayer)	[(None, 200)]	0	[]
lstm_1 (LSTM)	[(None, None, 200), (None, 200), (None, 200)]	240800	['dropout_1[0][0]', 'input_3[0][0]', 'input_4[0][0]']
input_5 (InputLayer)	[(None, None, 200)]	0	[]
dense (Dense)	(None, None, 2506)	503706	['lstm_1[1][0]']
=====			
Total params: 995,106			
Trainable params: 995,106			
Non-trainable params: 0			

Decoder Inference Model summary			
Model: "model_5"			
Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[(None, None)]	0	[]
embedding_3 (Embedding)	(None, None, 100)	250600	['input_7[0][0]']
dropout_3 (Dropout)	(None, None, 100)	0	['embedding_3[0][0]']
input_8 (InputLayer)	[(None, 200)]	0	[]
input_9 (InputLayer)	[(None, 200)]	0	[]
input_10 (InputLayer)	[(None, None, 200)]	0	[]
lstm_3 (LSTM)	[(None, None, 200), (None, 200), (None, 200)]	240800	['dropout_3[0][0]', 'input_8[0][0]', 'input_9[0][0]']
attention_layer_1 (AttentionLayer)	(None, None, 400)	0	['input_10[0][0]', 'lstm_3[1][0]']
dense_1 (Dense)	(None, None, 2506)	1004906	['attention_layer_1[0][0]']
=====			
Total params: 1,496,306			
Trainable params: 1,496,306			
Non-trainable params: 0			

We use the decoder model to create the decoder states rather than the encoder states. We then feed the LSTM the list of encoder states. In the LSTM, the target embeddings are taken as input, and the starting state is used to build a decoder. Our LSTM produces three outputs, which are assigned to variable.

The attention layer condition will be added next, which is the same as the one in encoder. The LSTM or attention layer output will be transmitted to the last dense layer, which will assign probability for the following token. This test set had a BLEU score of 5.06

3. Adding attention

```

"""
Task 3 attention

Start
"""

luong_score = tf.matmul(decoder_outputs, encoder_outputs, transpose_b=True)
alignment = tf.nn.softmax(luong_score, axis=2)
context = tf.matmul(K.expand_dims(alignment,axis=2), K.expand_dims(encoder_outputs,axis=1))
encoder_vector = K.squeeze(context,axis=2)

"""
End Task 3
"""

```

Putting together the decoder states

Decoder Inference Model summary

Model: "model_5"

Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[(None, None)]	0	[]
embedding_3 (Embedding)	(None, None, 100)	250600	['input_7[0][0]']
dropout_3 (Dropout)	(None, None, 100)	0	['embedding_3[0][0]']
input_8 (InputLayer)	[(None, 200)]	0	[]
input_9 (InputLayer)	[(None, 200)]	0	[]
input_10 (InputLayer)	[(None, None, 200)]	0	[]
lstm_3 (LSTM)	[(None, None, 200), (None, 200), (None, 200)]	240800	['dropout_3[0][0]', 'input_8[0][0]', 'input_9[0][0]']
attention_layer_1 (AttentionLayer)	(None, None, 400)	0	['input_10[0][0]', 'lstm_3[1][0]']
dense_1 (Dense)	(None, None, 2506)	1004906	['attention_layer_1[0][0]']

=====
Total params: 1,496,306
Trainable params: 1,496,306
Non-trainable params: 0

We'll now use the NMT's attention layer to boost our BLEU score. To begin, we must determine the "luong score," which is obtained by multiplying the decoder and encoder outputs in a matrix. We must transpose the "encoder outputs" since the shapes do not match the decoder output. Instead of transposing and multiplying, we used tensor flow's "matmul" function to combine the two steps.

Once the score has been computed, we should softmax the dimension with the size "max source sent len". We are doing this for axis 2. This modification is performed using softmax in TensorFlow.

The next step is to increase the dimension of "encoder output" and our softmax output so that we can multiply them element by element. Using the encoder vector, we multiplied the dimensions once they were enlarged. To get the same dimension as before, we must conduct a sum because multiplication updates the dimension. Our encoder vector is derived by summing the "max source sent len" values. This test set had a BLEU score of 15.33